

In-Class Exercise

- Input: social network graph stored as two files
 - Nodes: (userID, name, hobbies)
 - Edges: (userID, friendID)
- Find all cricket players in “my” 3-hop neighborhood
- Pre-process: create adjacency list for each person, add this to node object

// Initially only start node “me” is *active*. It is not *reached* and will remain so.

```
map(nid n, nodeObject N) {
```

```
  if (n <> “me”) emit(n, N) // Do not process start node any more
```

```
  if (N.isActive()) {
```

```
    N.unsetActive() // De-activate node to avoid repeat processing
```

```
    for all nid m in N.adjacencyList do
```

```
      if (m <> “me”) emit(m, NULL)
```

```
  }
```

```
}
```

```
// Identify newly reached nodes and set them to reached and active.
```

```
reduce(nid m, [n1, n2,...]) {
```

```
  isReached = false; M = NULL
```

```
  for all n in [n1,n2,...] do
```

```
    if isNode(n) then // The node object was found: recover graph structure
```

```
      M = n
```

```
    else // This node was reached from an active node. Set as reached.
```

```
      isReached = true
```

```
// Set this node as active if it is reached for the first time
```

```
if (not M.wasReached()) {
```

```
  M.setReached(); M.setActive()
```

```
}
```

```
emit(nid m, node M)
```

```
}
```

Filter Condition

- The program above is called 3 times.
 - Final result: 3-hop friends = nodes with `wasReached() = true`
- Map-only parallel scan of this result to find the cricket players
- Avoid forwarding reached nodes in Mapper?
 - Ok: handle NULL node object in reduce call, do not emit it
 - But: reached nodes are scattered over output of different iterations

In-Class Exercise

- Input: same social network graph
 - Nodes annotated with userID, name, hobbies
 - Edges = friendship links
- Find all cricket players in “my” 3-hop neighborhood and
- *Return them sorted by friendship distance*

Solution 1

- Pre-process: create adjacency list for each person, add this to node object
- Remember in which iteration the node was *first* reached
 - Pass iteration counter into context (global constant)
 - In Reduce, set a new distance variable to the iteration counter in the “if not M.wasReached()” block
- Run a sort post-processing step with distance variable as key, where map eliminates nodes that were not reached

Solution 2

- Pre-process: create adjacency list for each person, add this to node object
- Do not emit reached nodes in Mapper
- Friends discovered in i-th iteration are in output of i-th job
- Just pick them up in order from the corresponding output directories

```

// Initially only start node "me" is active. It is not reached and will remain so.
map(nid n, nodeObject N) {
  if (n <> "me" AND !N.wasReached()) emit(n, N) // Do not process start node and reached nodes again

  if (N.isActive()) {
    N.unsetActive() // De-activate node to avoid repeat processing
    for all nid m in N.adjacencyList do
      if (m <> "me") emit(m, NULL)
    }
  }
}

// Identify newly reached nodes and set them to reached and active.
reduce(nid m, [n1, n2,...]) {
  isReached = false; M = NULL

  for all n in [n1,n2,...] do
    if isNode(n) then // The node object was found: recover graph structure
      M = n
    else // This node was reached from an active node
      isReached = true

  // If the node object is NULL, the node was reached before: ignore it
  if (M <> NULL) {
    if (isReached) { M.setReached(); M.setActive(); }

    emit(nid m, node M)
  }
}

```